

Symphonia Consensus — Gambling Diagnostic Study Configuration

Prepared for: Lukas

Date: 26 March 2026

Status: Proposal

What Symphonia Already Does

Symphonia is a structured consensus platform that collects expert responses — free-text claims, evidence, and uncertainties — and synthesises them into traceable findings with provenance. It runs multiple independent analytical perspectives (a committee of dimensionally-confined agents) over the responses, cross-reviews their outputs, and produces a narrative report with agreement levels, minority positions, and source attribution. The system already supports multi-round Delphi workflows, where each round's synthesis feeds into the next.

In short: experts answer questions, the platform finds where they agree, where they disagree, and why — then writes it up with full traceability back to who said what.

What We Add for the Gambling Diagnostic Study

Your study design calls for clinicians to rate 20–30 diagnostic statements on Likert scales (e.g. 1–5: Strongly Disagree → Strongly Agree), optionally with free-text justification. The platform currently expects free-text responses. We need four additions:

1. Likert-Scale Input Mode

A new `LikertResponse` model that accepts numeric scale ratings per item, alongside optional free-text reasoning. This sits next to the existing `ExpertResponse` — not a replacement. Clinicians see a clean form with grouped statements and radio/slider inputs per item.

Effort: Small. New model + input adapter. The downstream pipeline already handles structured data — we're adding a front door, not rebuilding the house.

2. Statement Grouping & Sections

The instrument's 20–30 items are organised into diagnostic dimensions (e.g. behavioural indicators, cognitive distortions, harm severity). We add section metadata to the study configuration so items display grouped, and results are reported per-section as well as overall.

Effort: Small. Configuration-level — a section map in the study definition. The pipeline already processes per-question; we add a grouping layer on top.

3. Statistical Aggregation Layer

Likert data needs numeric summary statistics that the current narrative synthesis doesn't produce: per-item means, standard deviations, distribution shapes, inter-rater agreement (e.g. Kendall's W), and consensus thresholds. We add a lightweight statistical module that runs before narrative generation and feeds its outputs into the synthesis as structured findings.

The existing narrative engine then writes up the results in prose — “Panellists reached strong consensus ($W = 0.82$) on items related to cognitive distortions, while harm-severity indicators showed significant divergence” — grounded in the actual numbers rather than text extraction alone.

Effort: Moderate. New statistical module (~days of work), but it plugs cleanly into the existing pipeline between response collection and synthesis.

4. Voice Input (Nice-to-Have)

Allow clinicians to dictate free-text justifications via speech-to-text instead of typing. This is a front-end convenience layer — transcribed text feeds into the existing pipeline unchanged.

Effort: Small, but independent of the core work. Can be added later without affecting the study design.

What Stays the Same

Capability	Status
Multi-round Delphi workflow	Already works
Committee-based multi-perspective analysis	Already works
Narrative report generation with citations	Already works
Minority position preservation	Already works
Provenance tracking (who said what)	Already works
Cost tracking & checkpointing	Already works

The platform's analytical engine — committee deliberation, narrative synthesis, provenance, minority reports — runs unchanged. We're adding a numeric front door and a statistical summary layer. The heavy machinery is already built.

Timeline

Item	Estimate
Likert input model + form UI	1–2 days
Section grouping configuration	1 day

Item	Estimate
Statistical aggregation module	2–3 days
Integration testing with sample instrument	1 day
Voice input (if desired)	1–2 days (independent)
Total (core)	~1 week

Next Step

If this looks right, we'll need your instrument — the list of diagnostic statements, their groupings, and the scale definitions. We configure the study from that and you're running.